

Trax : Ungraded Lecture Notebook

In this notebook you'll get to know about the Trax framework and learn about some of its basic building blocks.

Background

Why Trax and not TensorFlow or PyTorch?

TensorFlow and PyTorch are both extensive frameworks that can do almost anything in deep learning. They offer a lot of flexibility, but that often means verbosity of syntax and extra time to code.

Trax is much more concise. It runs on a TensorFlow backend but allows you to train models with 1 line commands. Trax also runs end to end, allowing you to get data, model and train all with a single terse statements. This means you can focus on learning, instead of spending hours on the idiosyncrasies of big framework implementation.

Why not Keras then? ¶

Keras is now part of Tensorflow itself from 2.0 onwards. Also, trax is good for implementing new state of the art algorithms like Transformers, Reformers, BERT because it is actively maintained by Google Brain Team for advanced deep learning tasks. It runs smoothly on CPUs, GPUs and TPUs as well with comparatively lesser modifications in code.

How to Code in Trax

Building models in Trax relies on 2 key concepts:- **layers** and **combinators**. Trax layers are simple objects that process data and perform computations. They can be chained together into composite layers using Trax combinators, allowing you to build layers and models of any complexity.

Trax, JAX, TensorFlow and Tensor2Tensor

You already know that Trax uses Tensorflow as a backend, but it also uses the JAX library to speed up computation too. You can view JAX as an enhanced and optimized version of numpy.

Watch out for assignments which import `import trax.fastmath.numpy as np` . If you see this line, remember that when calling `np` you are really calling Trax's version of numpy that is compatible with JAX.

As a result of this, where you used to encounter the type `numpy.ndarray` now you will find the type `jax.interpreters.xla.DeviceArray` .

Tensor2Tensor is another name you might have heard. It started as an end to end solution much like how Trax is designed, but it grew unwieldy and complicated. So you can view Trax as the new improved version that operates much faster and simpler.

Resources

- Trax source code can be found on Github: [Trax \(https://github.com/google/trax\)](https://github.com/google/trax)
- JAX library: [JAX \(https://jax.readthedocs.io/en/latest/index.html\)](https://jax.readthedocs.io/en/latest/index.html)

Installing Trax

Trax has dependencies on JAX and some libraries like JAX which are yet to be supported in [Windows](https://github.com/google/jax/blob/1bc5896ee4eab5d7bb4ec6f161d8b2abb30557be/README.md#install) (<https://github.com/google/jax/blob/1bc5896ee4eab5d7bb4ec6f161d8b2abb30557be/README.md#install>) but work well in Ubuntu and MacOS. We would suggest that if you are working on Windows, try to install Trax on WSL2.

Official maintained documentation - [trax-ml](https://trax-ml.readthedocs.io/en/latest/) (<https://trax-ml.readthedocs.io/en/latest/>) not to be confused with this [TraX](https://trax.readthedocs.io/en/latest/index.html) (<https://trax.readthedocs.io/en/latest/index.html>)

```
In [12]: #!pip install trax==1.3.1 Use this version for this notebook
```

Imports

```
In [13]: import numpy as np # regular ol' numpy

from trax import layers as tl # core building block
from trax import shapes # data signatures: dimensionality and type
from trax import fastmath # uses jax, offers numpy on steroids
```

```
In [14]: # Trax version 1.3.1 or better
!pip list | grep trax
```

```
trax 1.3.1
WARNING: You are using pip version 20.1.1; however, version 20.2.2
is available.
You should consider upgrading via the '/opt/conda/bin/python -m pi
p install --upgrade pip' command.
```

Layers

Layers are the core building blocks in Trax or as mentioned in the lectures, they are the base classes.

They take inputs, compute functions/custom calculations and return outputs.

You can also inspect layer properties. Let me show you some examples.

Relu Layer

First I'll show you how to build a relu activation function as a layer. A layer like this is one of the simplest types. Notice there is no object initialization so it works just like a math function.

Note: Activation functions are also layers in Trax, which might look odd if you have been using other frameworks for a longer time.

```
In [15]: # Layers
# Create a relu trax layer
relu = tl.Relu()

# Inspect properties
print("-- Properties --")
print("name :", relu.name)
print("expected inputs :", relu.n_in)
print("promised outputs :", relu.n_out, "\n")

# Inputs
x = np.array([-2, -1, 0, 1, 2])
print("-- Inputs --")
print("x :", x, "\n")

# Outputs
y = relu(x)
print("-- Outputs --")
print("y :", y)
```

```
-- Properties --
name : Relu
expected inputs : 1
promised outputs : 1
```

```
-- Inputs --
x : [-2 -1  0  1  2]
```

```
-- Outputs --
y : [0 0 0 1 2]
```

Concatenate Layer

Now I'll show you how to build a layer that takes 2 inputs. Notice the change in the expected inputs property from 1 to 2.

```
In [16]: # Create a concatenate trax layer
concat = tl.Concatenate()
print("-- Properties --")
print("name :", concat.name)
print("expected inputs :", concat.n_in)
print("promised outputs :", concat.n_out, "\n")

# Inputs
x1 = np.array([-10, -20, -30])
x2 = x1 / -10
print("-- Inputs --")
print("x1 :", x1)
print("x2 :", x2, "\n")

# Outputs
y = concat([x1, x2])
print("-- Outputs --")
print("y :", y)
```

```
-- Properties --
name : Concatenate
expected inputs : 2
promised outputs : 1

-- Inputs --
x1 : [-10 -20 -30]
x2 : [1. 2. 3.]

-- Outputs --
y : [-10. -20. -30.   1.   2.   3.]
```

Layers are Configurable

You can change the default settings of layers. For example, you can change the expected inputs for a concatenate layer from 2 to 3 using the optional parameter `n_items`.

```
In [17]: # Configure a concatenate layer
concat_3 = tl.Concatenate(n_items=3) # configure the layer's expected inputs
print("-- Properties --")
print("name :", concat_3.name)
print("expected inputs :", concat_3.n_in)
print("promised outputs :", concat_3.n_out, "\n")

# Inputs
x1 = np.array([-10, -20, -30])
x2 = x1 / -10
x3 = x2 * 0.99
print("-- Inputs --")
print("x1 :", x1)
print("x2 :", x2)
print("x3 :", x3, "\n")

# Outputs
y = concat_3([x1, x2, x3])
print("-- Outputs --")
print("y :", y)
```

```
-- Properties --
name : Concatenate
expected inputs : 3
promised outputs : 1

-- Inputs --
x1 : [-10 -20 -30]
x2 : [1.  2.  3.]
x3 : [0.99 1.98 2.97]

-- Outputs --
y : [-10.   -20.   -30.    1.    2.    3.    0.99   1.98   2.97]
```

Note: At any point, if you want to refer the function help/ look up the [documentation \(https://trax-ml.readthedocs.io/en/latest/\)](https://trax-ml.readthedocs.io/en/latest/) or use help function.

```
In [18]: #help(tl.Concatenate) #Uncomment this to see the function docstring
         with explanation
```

Layers can have Weights

Some layer types include mutable weights and biases that are used in computation and training. Layers of this type require initialization before use.

For example the `LayerNorm` layer calculates normalized data, that is also scaled by weights and biases. During initialization you pass the data shape and data type of the inputs, so the layer can initialize compatible arrays of weights and biases.

```
In [19]: # Uncomment any of them to see information regarding the function
help(tl.LayerNorm)
# help(shapes.signature)
```

Help on class `LayerNorm` in module `trax.layers.normalization`:

```
class LayerNorm(trax.layers.base.Layer)
|   LayerNorm(epsilon=1e-06)
|
|   Layer normalization.
|
|   Method resolution order:
|       LayerNorm
|       trax.layers.base.Layer
|       builtins.object
|
|   Methods defined here:
|
|   __init__(self, epsilon=1e-06)
|       Creates a partially initialized, unconnected layer instance.
|
|   Args:
|       n_in: Number of inputs expected by this layer.
|       n_out: Number of outputs promised by this layer.
|       name: Class-like name for this layer; for use when printing this layer.
|       sublayers_to_print: Sublayers to display when printing out this layer;
|           By default (when None) we display all sublayers.
|
|   forward(self, x)
|       Computes this layer's output as part of a forward pass through the model.
|
|       Authors of new layer subclasses should override this method to define the
|       forward computation that their layer performs. Use `self.weights` to access
|       trainable weights of this layer. If you need to use local non-trainable
```

```

    |         state or randomness, use `self.rng` for the random seed (no
    |         need to set it)
    |         and use `self.state` for non-trainable state (and set it to
    |         the new value).
    |
    |         Args:
    |             inputs: Zero or more input tensors, packaged as described
    |             in the `Layer`
    |                 class docstring.
    |
    |         Returns:
    |             Zero or more output tensors, packaged as described in the
    |             `Layer` class
    |                 docstring.
    |
    |         init_weights_and_state(self, input_signature)
    |             Initializes weights and state for inputs with the given signature.
    |
    |         Authors of new layer subclasses should override this method if their layer
    |         uses trainable weights or non-trainable state. To initialize trainable
    |         weights, set `self.weights` and to initialize non-trainable state,
    |         set `self.state` to the intended value.
    |
    |         Args:
    |             input_signature: A `ShapeDtype` instance (if this layer
    |             takes one input)
    |                 or a list/tuple of `ShapeDtype` instances; signatures of inputs.
    |
    |         -----
    |         -----
    |
    |         Methods inherited from trax.layers.base.Layer:
    |
    |         __call__(self, x, weights=None, state=None, rng=None)
    |             Makes layers callable; for use in tests or interactive settings.
    |
    |             This convenience method helps library users play with, test, or otherwise
    |             probe the behavior of layers outside of a full training environment. It
    |             presents the layer as callable function from inputs to outputs, with the
    |             option of manually specifying weights and non-parameter state per individual
    |             call. For convenience, weights and non-parameter state are cached per layer
    |             instance, starting from default values of `EMPTY_WEIGHTS` and `EMPTY_STATE`,

```



```

|         and acquiring non-empty values either by initialization or
from values
|         explicitly provided via the weights and state keyword argu
ments.
|
|         Args:
|             x: Zero or more input tensors, packaged as described in
the `Layer` class
|             docstring.
|             weights: Weights or `None`; if `None`, use self's cached
weights value.
|             state: State or `None`; if `None`, use self's cached sta
te value.
|             rng: Single-use random number generator (JAX PRNG key),
or `None`;
|                 if `None`, use a default computed from an integer 0
seed.
|
|         Returns:
|             Zero or more output tensors, packaged as described in th
e `Layer` class
|             docstring.
|
|         __repr__(self)
|             Return repr(self).
|
|         backward(self, inputs, output, grad, weights, state, new_state
, rng)
|             Custom backward pass to propagate gradients in a custom wa
y.
|
|         Args:
|             inputs: Input tensors; can be a (possibly nested) tuple.
|             output: The result of running this layer on inputs.
|             grad: Gradient signal computed based on subsequent layer
s; its structure
|                 and shape must match output.
|             weights: This layer's weights.
|             state: This layer's state prior to the current forward p
ass.
|             new_state: This layer's state after the current forward
pass.
|             rng: Single-use random number generator (JAX PRNG key).
|
|         Returns:
|             The custom gradient signal for the input. Note that we n
eed to return
|                 a gradient for each argument of forward, so it will usua
lly be a tuple
|                 of signals: the gradient for inputs and weights.
|
|         init(self, input_signature, rng=None, use_cache=False)
|             Initializes weights/state of this layer and its sublayers

```

```

recursively.
|
|     Initialization creates layer weights and state, for layers
that use them.
|     It derives the necessary array shapes and data types from
the layer's input
|     signature, which is itself just shape and data type inform
ation.
|
|     For layers without weights or state, this method safely do
es nothing.
|
|     This method is designed to create weights/state only once
for each layer
|     instance, even if the same layer instance occurs in multip
le places in the
|     network. This enables weight sharing to be implemented as
layer sharing.
|
|     Args:
|         input_signature: `ShapeDtype` instance (if this layer ta
kes one input)
|         or list/tuple of `ShapeDtype` instances.
|         rng: Single-use random number generator (JAX PRNG key),
or `None`;
|         if `None`, use a default computed from an integer 0
seed.
|         use_cache: If `True`, and if this layer instance has alr
eady been
|         initialized elsewhere in the network, then return sp
ecial marker
|         values -- tuple `(GET_WEIGHTS_FROM_CACHE, GET_STATE_
FROM_CACHE)`.
|         Else return this layer's newly initialized weights a
nd state.
|
|     Returns:
|         A `(weights, state)` tuple.
|
|     init_from_file(self, file_name, weights_only=False, input_sign
ature=None)
|     Initializes this layer and its sublayers from a pickled ch
eckpoint.
|
|     In the common case (`weights_only=False`), the file must b
e a gzipped pickled
|     dictionary containing items with keys `'flat_weights', 'f
lat_state'` and
|     `'input_signature'`, which are used to initialize this lay
er.
|     If `input_signature` is specified, it's used instead of th
e one in the file.
|     If `weights_only` is `True`, the dictionary does not need

```

```

to have the
    |         ``flat_state`` item and the state it not restored either.
    |
    |         Args:
    |             file_name: Name/path of the pickled weights/state file.
    |             weights_only: If ``True``, initialize only the layer's wei
ghts. Else
    |                 initialize both weights and state.
    |             input_signature: Input signature to be used instead of t
he one from file.
    |
    |         output_signature(self, input_signature)
    |             Returns output signature this layer would give for ``input_
signature``.
    |
    |         pure_fn(self, x, weights, state, rng, use_cache=False)
    |             Applies this layer as a pure function with no optional arg
s.
    |
    |         This method exposes the layer's computation as a pure func
tion. This is
    |             especially useful for JIT compilation. Do not override, us
e ``forward``
    |             instead.
    |
    |         Args:
    |             x: Zero or more input tensors, packaged as described in
the ``Layer`` class
    |                 docstring.
    |             weights: A tuple or list of trainable weights, with one
element for this
    |                 layer if this layer has no sublayers, or one for eac
h sublayer if
    |                 this layer has sublayers. If a layer (or sublayer) h
as no trainable
    |                 weights, the corresponding weights element is an emp
ty tuple.
    |             state: Layer-specific non-parameter state that can updat
e between batches.
    |             rng: Single-use random number generator (JAX PRNG key).
    |             use_cache: if ``True``, cache weights and state in the lay
er object; used
    |                 to implement layer sharing in combinators.
    |
    |         Returns:
    |             A tuple of ``(tensors, state)``. The tensors match the num
ber (``n_out``)
    |             promised by this layer, and are packaged as described in
the ``Layer``
    |                 class docstring.
    |
    |         weights_and_state_signature(self, input_signature)
    |             Return a pair containing the signatures of weights and sta

```

te.

```

-----
Data descriptors inherited from trax.layers.base.Layer:
    __dict__
        dictionary for instance variables (if defined)
    __weakref__
        list of weak references to the object (if defined)
    has_backward
        Returns `True` if this layer provides its own custom backward
        and pass code.

        A layer subclass that provides custom backward pass code (
        for custom
        gradients) must override this method to return `True`.
    n_in
        Returns how many tensors this layer expects as input.
    n_out
        Returns how many tensors this layer promises as output.
    name
        Returns the name of this layer.
    rng
        Returns a single-use random number generator without advancing it.
    state
        Returns a tuple containing this layer's state; may be empty.
    sublayers
        Returns a tuple containing this layer's sublayers; may be empty.
    weights
        Returns this layer's weights.

        Depending on the layer, the weights can be in the form of:

        - an empty tuple
        - a tensor (ndarray)
        - a nested structure of tuples and tensors

```

```
In [20]: # Layer initialization
norm = tl.LayerNorm()
# You first must know what the input data will look like
x = np.array([0, 1, 2, 3], dtype="float")

# Use the input data signature to get shape and type for initializing weights and biases
norm.init(shapes.signature(x)) # We need to convert the input datatype from usual tuple to trax ShapeDtype

print("Normal shape:", x.shape, "Data Type:", type(x.shape))
print("Shapes Trax:", shapes.signature(x), "Data Type:", type(shapes.signature(x)))

# Inspect properties
print("-- Properties --")
print("name :", norm.name)
print("expected inputs :", norm.n_in)
print("promised outputs :", norm.n_out)
# Weights and biases
print("weights :", norm.weights[0])
print("biases :", norm.weights[1], "\n")

# Inputs
print("-- Inputs --")
print("x :", x)

# Outputs
y = norm(x)
print("-- Outputs --")
print("y :", y)
```

```
Normal shape: (4,) Data Type: <class 'tuple'>
Shapes Trax: ShapeDtype{shape:(4,), dtype:float64} Data Type: <class 'trax.shapes.ShapeDtype'>
-- Properties --
name : LayerNorm
expected inputs : 1
promised outputs : 1
weights : [1. 1. 1. 1.]
biases : [0. 0. 0. 0.]

-- Inputs --
x : [0. 1. 2. 3.]
-- Outputs --
y : [-1.3416404 -0.44721344  0.44721344  1.3416404 ]
```

Custom Layers

This is where things start getting more interesting! You can create your own custom layers too and define custom functions for computations by using `tl.Fn`. Let me show you how.

In [21]: `help(tl.Fn)`

Help on function Fn in module trax.layers.base:

`Fn(name, f, n_out=1)`

Returns a layer with no weights that applies the function `f`.

`f` can take and return any number of arguments, and takes only positional

arguments -- no default or keyword arguments. It often uses JAX-numpy (`jnp`).

The following, for example, would create a layer that takes two inputs and

returns two outputs -- element-wise sums and maxima:

```
tl.Fn('SumAndMax', lambda x0, x1: (x0 + x1, jnp.maximum(x0, x1)), n_out=2)
```

The layer's number of inputs (`n_in`) is automatically set to number of

positional arguments in `f`, but you must explicitly set the number of

outputs (`n_out`) whenever it's not the default value 1.

Args:

`name`: Class-like name for the resulting layer; for use in debugging.

`f`: Pure function from input tensors to output tensors, where each input

tensor is a separate positional arg, e.g., `f(x0, x1) --> x0 + x1`.

Output tensors must be packaged as specified in the `Layer` class

`docstring`.

`n_out`: Number of outputs promised by the layer; default value 1.

Returns:

Layer executing the function `f`.

```
In [22]: # Define a custom layer
# In this example you will create a layer to calculate the input times 2

def TimesTwo():
    layer_name = "TimesTwo" #don't forget to give your custom layer a name to identify

    # Custom function for the custom layer
    def func(x):
        return x * 2

    return tl.Fn(layer_name, func)

# Test it
times_two = TimesTwo()

# Inspect properties
print("-- Properties --")
print("name :", times_two.name)
print("expected inputs :", times_two.n_in)
print("promised outputs :", times_two.n_out, "\n")

# Inputs
x = np.array([1, 2, 3])
print("-- Inputs --")
print("x :", x, "\n")

# Outputs
y = times_two(x)
print("-- Outputs --")
print("y :", y)
```

```
-- Properties --
name : TimesTwo
expected inputs : 1
promised outputs : 1
```

```
-- Inputs --
x : [1 2 3]
```

```
-- Outputs --
y : [2 4 6]
```

Combinators

You can combine layers to build more complex layers. Trax provides a set of objects named combinator layers to make this happen. Combinators are themselves layers, so behavior commutes.

Serial Combinator

This is the most common and easiest to use. For example could build a simple neural network by combining layers into a single layer using the `Serial` combinator. This new layer then acts just like a single layer, so you can inspect inputs, outputs and weights. Or even combine it into another layer! Combinators can then be used as trainable models. *Try adding more layers*

Note:As you must have guessed, if there is serial combinator, there must be a parallel combinator as well. Do try to explore about combinators and other layers from the trax documentation and look at the repo to understand how these layers are written.

```
In [24]: # help(tl.Serial)
         # help(tl.Parallel)
```



```

In [25]: # Serial combinator
serial = tl.Serial(
    tl.LayerNorm(),          # normalize input
    tl.Relu(),               # convert negative values to zero
    times_two,               # the custom layer you created above, m
                             # multiplies the input recieved from above by 2

    ### START CODE HERE
    #     tl.Dense(n_units=2), # try adding more layers. eg uncomment
    #     these lines
    #     tl.Dense(n_units=1), # Binary classification, maybe? uncommen
    #     nt at your own peril
    #     tl.LogSoftmax()      # Yes, LogSoftmax is also a layer
    ### END CODE HERE
)

# Initialization
x = np.array([-2, -1, 0, 1, 2]) #input
serial.init(shapes.signature(x)) #initialising serial instance

print("-- Serial Model --")
print(serial, "\n")
print("-- Properties --")
print("name :", serial.name)
print("sublayers :", serial.sublayers)
print("expected inputs :", serial.n_in)
print("promised outputs :", serial.n_out)
print("weights & biases:", serial.weights, "\n")

# Inputs
print("-- Inputs --")
print("x :", x, "\n")

# Outputs
y = serial(x)
print("-- Outputs --")
print("y :", y)

```

```

-- Serial Model --
Serial[
  LayerNorm
  Relu
  TimesTwo
]

-- Properties --
name : Serial
sublayers : [LayerNorm, Relu, TimesTwo]
expected inputs : 1
promised outputs : 1
weights & biases: [(DeviceArray([1, 1, 1, 1, 1], dtype=int32), DeviceArray([0, 0, 0, 0, 0], dtype=int32)), (), ()]

-- Inputs --
x : [-2 -1  0  1  2]

-- Outputs --
y : [0.          0.          0.          1.4142132  2.8284264]

```

JAX

Just remember to lookout for which numpy you are using, the regular ol' numpy or Trax's JAX compatible numpy. Both tend to use the alias np so watch those import blocks.

Note: There are certain things which are still not possible in fastmath.numpy which can be done in numpy so you will see in assignments we will switch between them to get our work done.

```

In [26]: # Numpy vs fastmath.numpy have different data types
# Regular ol' numpy
x_numpy = np.array([1, 2, 3])
print("good old numpy : ", type(x_numpy), "\n")

# Fastmath and jax numpy
x_jax = fastmath.numpy.array([1, 2, 3])
print("jax trax numpy : ", type(x_jax))

good old numpy :  <class 'numpy.ndarray'>

jax trax numpy :  <class 'jax.interpreters.xla.DeviceArray'>

```

Summary

Trax is a concise framework, built on TensorFlow, for end to end machine learning. The key building blocks are layers and combinators. This notebook is just a taste, but sets you up with some key intuitions to take forward into the rest of the course and assignments where you will build end to end models.