

Assignment 4: Question duplicates

Welcome to the fourth assignment of course 3. In this assignment you will explore Siamese networks applied to natural language processing. You will further explore the fundamentals of Trax and you will be able to implement a more complicated structure using it. By completing this assignment, you will learn how to implement models with different architectures.

Outline

- [Overview](#)
- [Part 1: Importing the Data](#)
 - [1.1 Loading in the data](#)
 - [1.2 Converting a question to a tensor](#)
 - [1.3 Understanding the iterator](#)
 - [Exercise 01](#)
- [Part 2: Defining the Siamese model](#)
 - [2.1 Understanding Siamese Network](#)
 - [Exercise 02](#)
 - [2.2 Hard Negative Mining](#)
 - [Exercise 03](#)
- [Part 3: Training](#)
 - [3.1 Training the model](#)
 - [Exercise 04](#)
- [Part 4: Evaluation](#)
 - [4.1 Evaluating your siamese network](#)
 - [4.2 Classify](#)
 - [Exercise 05](#)
- [Part 5: Testing with your own questions](#)
 - [Exercise 06](#)
- [On Siamese networks](#)

Overview

In this assignment, concretely you will:

- Learn about Siamese networks
- Understand how the triplet loss works
- Understand how to evaluate accuracy
- Use cosine similarity between the model's outputted vectors
- Use the data generator to get batches of questions
- Predict using your own model

By now, you are familiar with trax and know how to make use of classes to define your model. We will start this homework by asking you to preprocess the data the same way you did in the previous assignments. After processing the data you will build a classifier that will allow you to identify whether two questions are the same or not.



You will process the data first and then pad in a similar way you have done in the previous assignment. Your model will take in the two question embeddings, run them through an LSTM, and then compare the outputs of the two sub networks using cosine similarity. Before taking a deep dive into the model, start by importing the data set.

Part 1: Importing the Data

1.1 Loading in the data

You will be using the Quora question answer dataset to build a model that could identify similar questions. This is a useful task because you don't want to have several versions of the same question posted. Several times when teaching I end up responding to similar questions on piazza, or on other community forums. This data set has been labeled for you. Run the cell below to import some of the packages you will be using.

```
In [1]: import os
import nltk
import trax
from trax import layers as tl
from trax.supervised import training
from trax.fastmath import numpy as fastnp
import numpy as np
import pandas as pd
import random as rnd

# set random seeds
trax.supervised.trainer_lib.init_random_number_generators(34)
rnd.seed(34)
```

INFO:tensorflow:tokens_length=568 inputs_length=512 targets_length=114 noise_density=0.15 mean_noise_span_length=3.0

Notice that for this assignment Trax's numpy is referred to as `fastnp` , while regular numpy is referred to as `np` .

You will now load in the data set. We have done some preprocessing for you. If you have taken the deeplearning specialization, this is a slightly different training method than the one you have seen there. If you have not, then don't worry about it, we will explain everything.

```
In [2]: data = pd.read_csv("questions.csv")
N=len(data)
print('Number of question pairs: ', N)
data.head()
```

Number of question pairs: 404351

Out[2]:

	id	qid1	qid2	question1	question2	is_duplicate
0	0	1	2	What is the step by step guide to invest in sh...	What is the step by step guide to invest in sh...	0
1	1	3	4	What is the story of Kohinoor (Koh-i-Noor) Dia...	What would happen if the Indian government sto...	0
2	2	5	6	How can I increase the speed of my internet co...	How can Internet speed be increased by hacking...	0
3	3	7	8	Why am I mentally very lonely? How can I solve...	Find the remainder when 23^{24} i...	0
4	4	9	10	Which one dissolve in water quikly sugar, salt...	Which fish would survive in salt water?	0

We first split the data into a train and test set. The test set will be used later to evaluate our model.

```
In [3]: N_train = 300000
N_test  = 10*1024
data_train = data[:N_train]
data_test  = data[N_train:N_train+N_test]
print("Train set:", len(data_train), "Test set:", len(data_test))
del(data) # remove to free memory
```

Train set: 300000 Test set: 10240

As explained in the lectures, we select only the question pairs that are duplicate to train the model. We build two batches as input for the Siamese network and we assume that question $q1_i$ (question i in the first batch) is a duplicate of $q2_i$ (question i in the second batch), but all other questions in the second batch are not duplicates of $q1_i$. The test set uses the original pairs of questions and the status describing if the questions are duplicates.

```
In [4]: td_index = (data_train['is_duplicate'] == 1).to_numpy()
td_index = [i for i, x in enumerate(td_index) if x]
print('number of duplicate questions: ', len(td_index))
print('indexes of first ten duplicate questions:', td_index[:10])
```

number of duplicate questions: 111486
indexes of first ten duplicate questions: [5, 7, 11, 12, 13, 15, 16, 18, 20, 29]

```
In [5]: print(data_train['question1'][5]) # Example of question duplicates (first one in data)
print(data_train['question2'][5])
print('is_duplicate: ', data_train['is_duplicate'][5])
```

Astrology: I am a Capricorn Sun Cap moon and cap rising...what does that say about me?
I'm a triple Capricorn (Sun, Moon and ascendant in Capricorn) What does this say about me?
is_duplicate: 1

```
In [6]: Q1_train_words = np.array(data_train['question1'][td_index])
Q2_train_words = np.array(data_train['question2'][td_index])

Q1_test_words = np.array(data_test['question1'])
Q2_test_words = np.array(data_test['question2'])
y_test = np.array(data_test['is_duplicate'])
```

Above, you have seen that you only took the duplicated questions for training our model. You did so on purpose, because the data generator will produce batches $([q_{1_1}, q_{1_2}, q_{1_3}, \dots], [q_{2_1}, q_{2_2}, q_{2_3}, \dots])$ where q_{1_i} and q_{2_k} are duplicate if and only if $i = k$.

Let's print to see what your data looks like.

```
In [7]: print('TRAINING QUESTIONS:\n')
print('Question 1: ', Q1_train_words[0])
print('Question 2: ', Q2_train_words[0], '\n')
print('Question 1: ', Q1_train_words[5])
print('Question 2: ', Q2_train_words[5], '\n')

print('TESTING QUESTIONS:\n')
print('Question 1: ', Q1_test_words[0])
print('Question 2: ', Q2_test_words[0], '\n')
print('is_duplicate =', y_test[0], '\n')
```

TRAINING QUESTIONS:

Question 1: Astrology: I am a Capricorn Sun Cap moon and cap rising...what does that say about me?

Question 2: I'm a triple Capricorn (Sun, Moon and ascendant in Capricorn) What does this say about me?

Question 1: What would a Trump presidency mean for current international master's students on an F1 visa?

Question 2: How will a Trump presidency affect the students presently in US or planning to study in US?

TESTING QUESTIONS:

Question 1: How do I prepare for interviews for cse?

Question 2: What is the best way to prepare for cse?

is_duplicate = 0

You will now encode each word of the selected duplicate pairs with an index. Given a question, you can then just encode it as a list of numbers.

First you tokenize the questions using `nltk.word_tokenize`.

You need a python default dictionary which later, during inference, assigns the values 0 to all Out Of Vocabulary (OOV) words.

Then you encode each word of the selected duplicate pairs with an index. Given a question, you can then just encode it as a list of numbers.

```
In [8]: #create arrays
Q1_train = np.empty_like(Q1_train_words)
Q2_train = np.empty_like(Q2_train_words)

Q1_test = np.empty_like(Q1_test_words)
Q2_test = np.empty_like(Q2_test_words)
```

```
In [9]: # Building the vocabulary with the train set           (this might take a minute)
from collections import defaultdict

vocab = defaultdict(lambda: 0)
vocab['<PAD>'] = 1

for idx in range(len(Q1_train_words)):
    Q1_train[idx] = nltk.word_tokenize(Q1_train_words[idx])
    Q2_train[idx] = nltk.word_tokenize(Q2_train_words[idx])
    q = Q1_train[idx] + Q2_train[idx]
    for word in q:
        if word not in vocab:
            vocab[word] = len(vocab) + 1
print('The length of the vocabulary is: ', len(vocab))
```

The length of the vocabulary is: 36268

```
In [10]: print(vocab['<PAD>'])
print(vocab['Astrology'])
print(vocab['Astronomy']) #not in vocabulary, returns 0
```

1
2
0

```
In [11]: for idx in range(len(Q1_test_words)):
    Q1_test[idx] = nltk.word_tokenize(Q1_test_words[idx])
    Q2_test[idx] = nltk.word_tokenize(Q2_test_words[idx])
```

```
In [12]: print('Train set has reduced to: ', len(Q1_train) )
print('Test set length: ', len(Q1_test) )
```

Train set has reduced to: 111486
Test set length: 10240

1.2 Converting a question to a tensor

You will now convert every question to a tensor, or an array of numbers, using your vocabulary built above.

```
In [13]: # Converting questions to array of integers
for i in range(len(Q1_train)):
    Q1_train[i] = [vocab[word] for word in Q1_train[i]]
    Q2_train[i] = [vocab[word] for word in Q2_train[i]]

for i in range(len(Q1_test)):
    Q1_test[i] = [vocab[word] for word in Q1_test[i]]
    Q2_test[i] = [vocab[word] for word in Q2_test[i]]
```

```
In [14]: print('first question in the train set:\n')
print(Q1_train_words[0], '\n')
print('encoded version:')
print(Q1_train[0], '\n')

print('first question in the test set:\n')
print(Q1_test_words[0], '\n')
print('encoded version:')
print(Q1_test[0])
```

first question in the train set:

Astrology: I am a Capricorn Sun Cap moon and cap rising...what does that say about me?

encoded version:

[2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21]

first question in the test set:

How do I prepare for interviews for cse?

encoded version:

[32, 38, 4, 107, 65, 1015, 65, 11509, 21]

You will now split your train set into a training/validation set so that you can use it to train and evaluate your Siamese model.

```
In [15]: # Splitting the data
cut_off = int(len(Q1_train)*.8)
train_Q1, train_Q2 = Q1_train[:cut_off], Q2_train[:cut_off]
val_Q1, val_Q2 = Q1_train[cut_off:], Q2_train[cut_off:]
print('Number of duplicate questions: ', len(Q1_train))
print("The length of the training set is: ", len(train_Q1))
print("The length of the validation set is: ", len(val_Q1))
```

Number of duplicate questions: 111486

The length of the training set is: 89188

The length of the validation set is: 22298

1.3 Understanding the iterator

Most of the time in Natural Language Processing, and AI in general we use batches when training our data sets. If you were to use stochastic gradient descent with one example at a time, it will take you forever to build a model. In this example, we show you how you can build a data generator that takes in $Q1$ and $Q2$ and returns a batch of size `batch_size` in the following format $([q1_1, q1_2, q1_3, \dots], [q2_1, q2_2, q2_3, \dots])$. The tuple consists of two arrays and each array has `batch_size` questions. Again, $q1_i$ and $q2_i$ are duplicates, but they are not duplicates with any other elements in the batch.

The command `next(data_generator)` returns the next batch. This iterator returns the data in a format that you could directly use in your model when computing the feed-forward of your algorithm. This iterator returns a pair of arrays of questions.

Exercise 01

Instructions:

Implement the data generator below. Here are some things you will need.

- While true loop.
- if `index >= len_Q1`, set the `idx` to 0.
- The generator should return shuffled batches of data. To achieve this without modifying the actual question lists, a list containing the indexes of the questions is created. This list can be shuffled and used to get random batches everytime the index is reset.
- Append elements of $Q1$ and $Q2$ to `input1` and `input2` respectively.
- if `len(input1) == batch_size`, determine `max_len` as the longest question in `input1` and `input2`. Ceil `max_len` to a power of 2 (for computation purposes) using the following command: `max_len = 2**int(np.ceil(np.log2(max_len)))`.
- Pad every question by `vocab['<PAD>']` until you get the length `max_len`.
- Use `yield` to return `input1, input2`.
- Don't forget to reset `input1, input2` to empty arrays at the end (data generator resumes from where it last left).

```
In [16]: # UNQ_C1 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# GRADED FUNCTION: data_generator
def data_generator(Q1, Q2, batch_size, pad=1, shuffle=True):
    """Generator function that yields batches of data

    Args:
        Q1 (list): List of transformed (to tensor) questions.
        Q2 (list): List of transformed (to tensor) questions.
        batch_size (int): Number of elements per batch.
        pad (int, optional): Pad character from the vocab. Defaults
        to 1.
```



```

        shuffle (bool, optional): If the batches should be randomi
zed or not. Defaults to True.
    Yields:
        tuple: Of the form (input1, input2) with types (numpy.ndarr
ay, numpy.ndarray)
        NOTE: input1: inputs to your model [q1a, q2a, q3a, ...] i.e
. (q1a,q1b) are duplicates
            input2: targets to your model [q1b, q2b,q3b, ...] i.e
. (q1a,q2i) i!=a are not duplicates
        """

    input1 = []
    input2 = []
    idx = 0
    len_q = len(Q1)
    question_indexes = [*range(len_q)]

    if shuffle:
        rnd.shuffle(question_indexes)

    ### START CODE HERE (Replace instances of 'None' with your code
) ###
    while True:
        if idx >= len_q:
            # if idx is greater than or equal to len_q, set idx acc
ordingly

            # (Hint: look at the instructions above)
            idx = 0
            # shuffle to get random batches if shuffle is set to Tr
ue

            if shuffle:
                rnd.shuffle(question_indexes)

            # get questions at the `question_indexes[idx]` position in
Q1 and Q2
            q1 = Q1[question_indexes[idx]]
            q2 = Q2[question_indexes[idx]]

            # increment idx by 1
            idx += 1
            # append q1
            input1.append(q1)
            # append q2
            input2.append(q2)
            if len(input1) == batch_size:
                # determine max_len as the longest question in input1 &
input 2

                # Hint: use the `max` function.
                # take max of input1 & input2 and then max out of the t
wo of them.
                max_len = max(max([len(q) for q in input1]),max([len(q)
for q in input2]))
                # pad to power-of-2 (Hint: look at the instructions abo

```

```

ve)

max_len = 2**int(np.ceil(np.log2(max_len)))
b1 = []
b2 = []
for q1, q2 in zip(input1, input2):
    # add [pad] to q1 until it reaches max_len
    q1 = q1 + [pad]*(max_len-len(q1))
    # add [pad] to q2 until it reaches max_len
    q2 = q2 + [pad]*(max_len-len(q2))
    # append q1
    b1.append(q1)
    # append q2
    b2.append(q2)
    # use b1 and b2
    yield np.array(b1), np.array(b2)
### END CODE HERE ###
# reset the batches
input1, input2 = [], [] # reset the batches

```

```

In [17]: batch_size = 2
res1, res2 = next(data_generator(train_Q1, train_Q2, batch_size))
print("First questions : ", '\n', res1, '\n')
print("Second questions : ", '\n', res2)

```

```

First questions :
[[ 30  87  78 134 2132 1981  28  78 594  21  1  1
 1  1
 1  1]
 [ 30  55  78 3541 1460  28  56 253  21  1  1  1  1
 1
 1  1]]

Second questions :
[[ 30 156  78 134 2132 9508  21  1  1  1  1  1
 1  1
 1  1]
 [ 30 156  78 3541 1460 131  56 253  21  1  1  1  1
 1
 1  1]]

```

Note: The following expected output is valid only if you run the above test cell **once** (first time). The output will change on each execution.

If you think your implementation is correct and it is not matching the output, make sure to restart the kernel and run all the cells from the top again.

Expected Output:

First questions :

```
[[ 30  87  78 134 2132 1981  28  78 594  21  1  1  1
1
    1  1]
[ 30  55  78 3541 1460  28  56 253  21  1  1  1  1
1  1]]
```

Second questions :

```
[[ 30 156  78 134 2132 9508  21  1  1  1  1  1
1
    1  1]
[ 30 156  78 3541 1460 131  56 253  21  1  1  1  1
1  1]]
```


Now that you have your generator, you can just call it and it will return tensors which correspond to your questions in the Quora data set.

Now you can go ahead and start building your neural network.

Part 2: Defining the Siamese model

2.1 Understanding Siamese Network

A Siamese network is a neural network which uses the same weights while working in tandem on two different input vectors to compute comparable output vectors. The Siamese network you are about to implement looks like this:



You get the question embedding, run it through an LSTM layer, normalize v_1 and v_2 , and finally use a triplet loss (explained below) to get the corresponding cosine similarity for each pair of questions. As usual, you will start by importing the data set. The triplet loss makes use of a baseline (anchor) input that is compared to a positive (truthy) input and a negative (falsy) input. The distance from the baseline (anchor) input to the positive (truthy) input is minimized, and the distance from the baseline (anchor) input to the negative (falsy) input is maximized. In math equations, you are trying to maximize the following.

$$\mathcal{L}(A, P, N) = \max(\|f(A) - f(P)\|^2 - \|f(A) - f(N)\|^2 + \alpha, 0)$$

A is the anchor input, for example $q1_1$, P the duplicate input, for example, $q2_1$, and N the negative input (the non duplicate question), for example $q2_2$.

α is a margin; you can think about it as a safety net, or by how much you want to push the duplicates from the non duplicates.

Exercise 02

Instructions: Implement the `Siamese` function below. You should be using all the objects explained below.

To implement this model, you will be using `trax`. Concretely, you will be using the following functions.

- `tl.Serial`: Combinator that applies layers serially (by function composition) allows you set up the overall structure of the feedforward. [docs \(https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.combinators.Serial\)](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.combinators.Serial) / [source code \(https://github.com/google/trax/blob/1372b903bb66b0daccee19fd0b1fdf44f659330b/trax/layers/combinators.py#L137\)](https://github.com/google/trax/blob/1372b903bb66b0daccee19fd0b1fdf44f659330b/trax/layers/combinators.py#L137)
 - You can pass in the layers as arguments to `Serial`, separated by commas.
 - For example: `tl.Serial(tl.Embeddings(...), tl.Mean(...), tl.Dense(...), tl.LogSoftmax(...))`
- `tl.Embedding`: Maps discrete tokens to vectors. It will have shape (vocabulary length X dimension of output vectors). The dimension of output vectors (also called `d_feature`) is the number of elements in the word embedding. [docs \(https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.embedding.Embedding\)](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.embedding.Embedding)

ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.core.Embedding) / [source code](#)
<https://github.com/google/trax/blob/1372b903bb66b0daccee19fd0b1fdf44f659330b/trax/layers/core>.

- `tl.Embedding(vocab_size, d_feature)` .
 - `vocab_size` is the number of unique words in the given vocabulary.
 - `d_feature` is the number of elements in the word embedding (some choices for a word embedding size range from 150 to 300, for example).
- `tl.LSTM` The LSTM layer. It leverages another Trax layer called `LSTMCell` (<https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.rnn.LSTMCell>). The number of units should be specified and should match the number of elements in the word embedding. [docs](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.rnn.LSTM) (<https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.rnn.LSTM>) / [source code](#)
<https://github.com/google/trax/blob/1372b903bb66b0daccee19fd0b1fdf44f659330b/trax/layers/rnn.r>
 - `tl.LSTM(n_units)` Builds an LSTM layer of `n_units`.
 - `tl.Mean` : Computes the mean across a desired axis. Mean uses one tensor axis to form groups of values and replaces each group with the mean value of that group. [docs](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.core.Mean) (<https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.core.Mean>) / [source code](#)
<https://github.com/google/trax/blob/1372b903bb66b0daccee19fd0b1fdf44f659330b/trax/layers/core>.
 - `tl.Mean(axis=1)` mean over columns.
 - `tl.Fn` Layer with no weights that applies the function `f`, which should be specified using a lambda syntax. [docs](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.base.Fn) (<https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.base.Fn>) / [source code](#)
<https://github.com/google/trax/blob/70f5364dc6f6ec11aabb918e5f5e4b0f5bfb995/trax/layers/base>.
 - `x ->` This is used for cosine similarity.
 - `tl.Fn('Normalize', lambda x: normalize(x))` Returns a layer with no weights that applies the function `f`
 - `tl.parallel` : It is a combinator layer (like `Serial`) that applies a list of layers in parallel to its inputs. [docs](https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.combinators.Parallel) (<https://trax-ml.readthedocs.io/en/latest/trax.layers.html#trax.layers.combinators.Parallel>) / [source code](#)
<https://github.com/google/trax/blob/37aba571a89a8ad86be76a569d0ec4a46bdd8642/trax/layers/co>

```

In [19]: # UNQ_C2 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# GRADED FUNCTION: Siamese
def Siamese(vocab_size=len(vocab), d_model=128, mode='train'):
    """Returns a Siamese model.

    Args:
        vocab_size (int, optional): Length of the vocabulary. Defaults to len(vocab).
        d_model (int, optional): Depth of the model. Defaults to 128.
        mode (str, optional): 'train', 'eval' or 'predict', predict mode is for fast inference. Defaults to 'train'.

    Returns:
        trax.layers.combinators.Parallel: A Siamese model.
    """

    def normalize(x): # normalizes the vectors to have L2 norm 1
        return x / fastnp.sqrt(fastnp.sum(x * x, axis=-1, keepdims=True))

    ### START CODE HERE (Replace instances of 'None' with your code) ###
    q_processor = tl.Serial( # Processor will run on Q1 and Q2.
        tl.Embedding(vocab_size=vocab_size, d_feature=d_model), # Embedding layer
        tl.LSTM(n_units=d_model), # LSTM layer
        tl.Mean(axis=1), # Mean over columns
        tl.Fn('Normalize', lambda x: normalize(x)) # Apply normalize function
    ) # Returns one vector of shape [batch_size, d_model].

    ### END CODE HERE ###

    # Run on Q1 and Q2 in parallel.
    model = tl.Parallel(q_processor, q_processor)
    return model

```

Setup the Siamese network model

```
In [20]: # check your model
model = Siamese()
print(model)

Parallel_in2_out2[
  Serial[
    Embedding_41699_128
    LSTM_128
    Mean
    Normalize
  ]
  Serial[
    Embedding_41699_128
    LSTM_128
    Mean
    Normalize
  ]
]
```

Expected output:

```
Parallel_in2_out2[
  Serial[
    Embedding_41699_128
    LSTM_128
    Mean
    Normalize
  ]
  Serial[
    Embedding_41699_128
    LSTM_128
    Mean
    Normalize
  ]
]
```

2.2 Hard Negative Mining

You will now implement the `TripletLoss`.

As explained in the lecture, loss is composed of two terms. One term utilizes the mean of all the non duplicates, the second utilizes the *closest negative*. Our loss expression is then:

$$\begin{aligned}\mathcal{L}_{loss_1}(\mathcal{A}, \mathcal{P}, \mathcal{N}) &= \max(-\cos(A, P) + \text{mean}_{neg} + \alpha, 0) \\ \mathcal{L}_{loss_2}(\mathcal{A}, \mathcal{P}, \mathcal{N}) &= \max(-\cos(A, P) + \text{closest}_{neg} + \alpha, 0) \\ \mathcal{L}_{loss}(\mathcal{A}, \mathcal{P}, \mathcal{N}) &= \text{mean}(\mathcal{L}_{loss_1} + \mathcal{L}_{loss_2})\end{aligned}$$

Further, two sets of instructions are provided. The first set provides a brief description of the task. If that set proves insufficient, a more detailed set can be displayed.

Exercise 03

Instructions (Brief): Here is a list of things you should do:

- As this will be run inside trax, use `fastnp.xyz` when using any `xyz` numpy function
- Use `fastnp.dot` to calculate the similarity matrix $v_1 v_2^T$ of dimension `batch_size` x `batch_size`
- Take the score of the duplicates on the diagonal `fastnp.diagonal`
- Use the trax functions `fastnp.eye` and `fastnp.maximum` for the identity matrix and the maximum.

► More Detailed Instructions

```
In [23]: # UNQ_C3 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# GRADED FUNCTION: TripletLossFn
def TripletLossFn(v1, v2, margin=0.25):
    """Custom Loss function.

    Args:
        v1 (numpy.ndarray): Array with dimension (batch_size, model
        _dimension) associated to Q1.
        v2 (numpy.ndarray): Array with dimension (batch_size, model
        _dimension) associated to Q2.
        margin (float, optional): Desired margin. Defaults to 0.25.

    Returns:
        jax.interpreters.xla.DeviceArray: Triplet Loss.
    """
    ### START CODE HERE (Replace instances of 'None' with your code)
    ###

    # use fastnp to take the dot product of the two batches (don't
    forget to transpose the second argument)
```



```

scores = fastnp.dot(v1, v2.T) # pairwise cosine sim
# calculate new batch size
batch_size = len(scores)
# use fastnp to grab all postive `diagonal` entries in `scores`
positive = fastnp.diagonal(scores) # the positive ones (duplicates)
# multiply `fastnp.eye(batch_size)` with 2.0 and subtract it out of `scores`
negative_without_positive = scores - fastnp.eye(batch_size)*2.0
# take the row by row `max` of `negative_without_positive`.
# Hint: negative_without_positive.max(axis = [?])
closest_negative = negative_without_positive.max(axis=1)
# subtract `fastnp.eye(batch_size)` out of 1.0 and do element-wise multiplication with `scores`
negative_zero_on_duplicate = scores*(1.0 - fastnp.eye(batch_size))
# use `fastnp.sum` on `negative_zero_on_duplicate` for `axis=1` and divide it by `(batch_size - 1)`
mean_negative = fastnp.sum(negative_zero_on_duplicate, axis=1)/(batch_size-1)
# compute `fastnp.maximum` among 0.0 and `A`
# A = subtract `positive` from `margin` and add `closest_negative`
triplet_loss1 = fastnp.maximum(0.0, margin - positive + closest_negative)
# compute `fastnp.maximum` among 0.0 and `B`
# B = subtract `positive` from `margin` and add `mean_negative`
triplet_loss2 = fastnp.maximum(0.0, margin - positive + mean_negative)
# add the two losses together and take the `fastnp.mean` of it
triplet_loss = fastnp.mean(triplet_loss1 + triplet_loss2)

### END CODE HERE ###

return triplet_loss

```

```

In [24]: v1 = np.array([[0.26726124, 0.53452248, 0.80178373],[0.5178918 , 0.57543534, 0.63297887]])
v2 = np.array([[ 0.26726124,  0.53452248,  0.80178373],[-0.5178918 , -0.57543534, -0.63297887]])
TripletLossFn(v2,v1)
print("Triplet Loss:", TripletLossFn(v2,v1))

```

Triplet Loss: 0.5

Expected Output:

Triplet Loss: 0.5

To make a layer out of a function with no trainable variables, use `tl.Fn`.

```
In [25]: from functools import partial
def TripletLoss(margin=0.25):
    triplet_loss_fn = partial(TripletLossFn, margin=margin)
    return tl.Fn('TripletLoss', triplet_loss_fn)
```

Part 3: Training

Now you are going to train your model. As usual, you have to define the cost function and the optimizer. You also have to feed in the built model. Before, going into the training, we will use a special data set up. We will define the inputs using the data generator we built above. The lambda function acts as a seed to remember the last batch that was given. Run the cell below to get the question pairs inputs.

```
In [26]: batch_size = 256
train_generator = data_generator(train_Q1, train_Q2, batch_size, vocab['<PAD>'])
val_generator = data_generator(val_Q1, val_Q2, batch_size, vocab['<PAD>'])
print('train_Q1.shape ', train_Q1.shape)
print('val_Q1.shape   ', val_Q1.shape)

train_Q1.shape (89188,)
val_Q1.shape   (22298,)
```

3.1 Training the model

You will now write a function that takes in your model and trains it. To train your model you have to decide how many times you want to iterate over the entire data set; each iteration is defined as an epoch . For each epoch, you have to go over all the data, using your training iterator.

Exercise 04

Instructions: Implement the `train_model` below to train the neural network above. Here is a list of things you should do, as already shown in lecture 7:

- Create `TrainTask` and `EvalTask`
- Create the training loop `trax.supervised.training.Loop`
- Pass in the following depending on the context (`train_task` or `eval_task`):
 - `labeled_data=generator`
 - `metrics=[TripletLoss()]`,
 - `loss_layer=TripletLoss()`
 - `optimizer=trax.optimizers.Adam` with learning rate of 0.01
 - `lr_schedule=lr_schedule`,
 - `output_dir=output_dir`

You will be using your triplet loss function with Adam optimizer. Please read the [trax \(https://trax-ml.readthedocs.io/en/latest/trax.optimizers.html?highlight=adam#trax.optimizers.adam.Adam\)](https://trax-ml.readthedocs.io/en/latest/trax.optimizers.html?highlight=adam#trax.optimizers.adam.Adam) documentation to get a full understanding.

This function should return a `training.Loop` object. To read more about this check the [docs \(https://trax-ml.readthedocs.io/en/latest/trax.supervised.html?highlight=loop#trax.supervised.training.Loop\)](https://trax-ml.readthedocs.io/en/latest/trax.supervised.html?highlight=loop#trax.supervised.training.Loop).

```
In [29]: lr_schedule = trax.lr.warmup_and_rsqrtd_decay(400, 0.01)

# UNQ_C4 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# GRADED FUNCTION: train_model
def train_model(Siamese, TripletLoss, lr_schedule, train_generator=
train_generator, val_generator=val_generator, output_dir='model/'):
    """Training the Siamese Model

    Args:
        Siamese (function): Function that returns the Siamese model
        .
        TripletLoss (function): Function that defines the TripletLoss loss function.
        lr_schedule (function): Trax multifactor schedule function.
        train_generator (generator, optional): Training generator.
        Defaults to train_generator.
        val_generator (generator, optional): Validation generator.
        Defaults to val_generator.
```

```

        output_dir (str, optional): Path to save model to. Defaults
to 'model/'.

Returns:
    trax.supervised.training.Loop: Training loop for the model.
    """
    output_dir = os.path.expanduser(output_dir)

    ### START CODE HERE (Replace instances of 'None' with your code
) ###

    train_task = training.TrainTask(
        labeled_data=train_generator,          # Use generator (train)
        loss_layer=TripletLoss(),             # Use triplet loss. Don't
forget to instantiate this object
        optimizer=trax.optimizers.Adam(learning_rate=0.01),
# Don't forget to add the learning rate parameter
        lr_schedule=lr_schedule, # Use Trax multifactor schedule fu
nction
    )

    eval_task = training.EvalTask(
        labeled_data=val_generator,           # Use generator (val)
        metrics=[TripletLoss()],             # Use triplet loss. Don't
forget to instantiate this object
    )

    ### END CODE HERE ###

    training_loop = training.Loop(Siamese(),
                                  train_task,
                                  eval_task=eval_task,
                                  output_dir=output_dir)

    return training_loop

```

```

In [30]: train_steps = 5
training_loop = train_model(Siamese, TripletLoss, lr_schedule)
training_loop.run(train_steps)

```

```

Step      1: train TripletLoss | 0.49954823
Step      1: eval  TripletLoss | 0.49950948

```

The model was only trained for 5 steps due to the constraints of this environment. For the rest of the assignment you will be using a pretrained model but now you should understand how the training can be done using Trax.

Part 4: Evaluation

4.1 Evaluating your siamese network

In this section you will learn how to evaluate a Siamese network. You will first start by loading a pretrained model and then you will use it to predict.

```
In [40]: # Loading in the saved model
model = Siamese()
model.init_from_file('model.pkl.gz')
```

4.2 Classify

To determine the accuracy of the model, we will utilize the test set that was configured earlier. While in training we used only positive examples, the test data, Q1_test, Q2_test and y_test, is setup as pairs of questions, some of which are duplicates some are not. This routine will run all the test question pairs through the model, compute the cosine similarity of each pair, threshold it and compare the result to y_test - the correct response from the data set. The results are accumulated to produce an accuracy.

Exercise 05

Instructions

- Loop through the incoming data in batch_size chunks
- Use the data generator to load q1, q2 a batch at a time. **Don't forget to set shuffle=False!**
- copy a batch_size chunk of y into y_test
- compute v1, v2 using the model
- for each element of the batch
 - compute the cos similarity of each pair of entries, v1[j],v2[j]
 - determine if d > threshold
 - increment accuracy if that result matches the expected results (y_test[j])
- compute the final accuracy and return

Due to some limitations of this environment, running classify multiple times may result in the kernel failing. If that happens *Restart Kernal & clear output* and then run from the top. During development, consider using a smaller set of data to reduce the number of calls to model().

```
In [41]: # UNQ_C5 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# GRADED FUNCTION: classify
```

```

def classify(test_Q1, test_Q2, y, threshold, model, vocab, data_gen
erator=data_generator, batch_size=64):
    """Function to test the accuracy of the model.

    Args:
        test_Q1 (numpy.ndarray): Array of Q1 questions.
        test_Q2 (numpy.ndarray): Array of Q2 questions.
        y (numpy.ndarray): Array of actual target.
        threshold (float): Desired threshold.
        model (trax.layers.combinators.Parallel): The Siamese model
        .
        vocab (collections.defaultdict): The vocabulary used.
        data_generator (function): Data generator function. Default
s to data_generator.
        batch_size (int, optional): Size of the batches. Defaults t
o 64.

    Returns:
        float: Accuracy of the model.
    """
    accuracy = 0
    ### START CODE HERE (Replace instances of 'None' with your code
) ###
    for i in range(0, len(test_Q1), batch_size):
        # Call the data generator (built in Ex 01) with shuffle=False using next()
        # use batch size chunks of questions as Q1 & Q2 arguments
        # of the data generator. e.g x[i:i + batch_size]
        # Hint: use `vocab['<PAD>']` for the `pad` argument of the
        data generator
        q1, q2 = next(data_generator(test_Q1[i:i + batch_size], tes
t_Q2[i:i + batch_size],
                                batch_size, vocab['<PAD>'], sh
uffle=False))
        # use batch size chunks of actual output targets (same syn
tax as example above)
        y_test = y[i:i + batch_size]
        # Call the model
        v1, v2 = model((q1, q2))

        for j in range(batch_size):
            # take dot product to compute cos similarity of each pa
            ir of entries, v1[j], v2[j]
            # don't forget to transpose the second argument
            d = np.dot(v1[j], v2[j].T)
            # is d greater than the threshold?
            res = d>threshold
            # increment accuracy if y_test is equal `res`
            accuracy += (y_test[j] == res)
        # compute accuracy using accuracy and total length of test ques
        tions
    accuracy = accuracy / len(test_Q1)
    ### END CODE HERE ###

```

```
return accuracy
```

```
In [42]: # this takes around 1 minute
accuracy = classify(Q1_test,Q2_test, y_test, 0.7, model, vocab, batch_size = 512)
print("Accuracy", accuracy)
```

Accuracy 0.69091796875

Expected Result

Accuracy ~0.69

Part 5: Testing with your own questions

In this section you will test the model with your own questions. You will write a function `predict` which takes two questions as input and returns 1 or 0 depending on whether the question pair is a duplicate or not.

But first, we build a reverse vocabulary that allows to map encoded questions back to words:

Write a function `predict` that takes in two questions, the model, and the vocabulary and returns whether the questions are duplicates (1) or not duplicates (0) given a similarity threshold.

Exercise 06

Instructions:

- Tokenize your question using `nltk.word_tokenize`
- Create Q1,Q2 by encoding your questions as a list of numbers using `vocab`
- pad Q1,Q2 with `next(data_generator([Q1], [Q2],1,vocab['']))`
- use `model()` to create `v1, v2`
- compute the cosine similarity (dot product) of `v1, v2`
- compute `res` by comparing `d` to the threshold

```
In [43]: # UNQ_C6 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# GRADED FUNCTION: predict
def predict(question1, question2, threshold, model, vocab, data_generator=data_generator, verbose=False):
    """Function for predicting if two questions are duplicates.

    Args:
        question1 (str): First question.
        question2 (str): Second question.
```

```

        threshold (float): Desired threshold.
        model (trax.layers.combinators.Parallel): The Siamese model
    .
        vocab (collections.defaultdict): The vocabulary used.
        data_generator (function): Data generator function. Defaults
s to data_generator.
        verbose (bool, optional): If the results should be printed
out. Defaults to False.

Returns:
    bool: True if the questions are duplicates, False otherwise
    .
    """
    ### START CODE HERE (Replace instances of 'None' with your code
) ###
    # use `nltk` word tokenize function to tokenize
    q1 = nltk.word_tokenize(question1) # tokenize
    q2 = nltk.word_tokenize(question2) # tokenize
    Q1, Q2 = [], []
    for word in q1: # encode q1
        # increment by checking the 'word' index in `vocab`
        Q1 += [vocab[word]]
    for word in q2: # encode q2
        # increment by checking the 'word' index in `vocab`
        Q2 += [vocab[word]]

    # Call the data generator (built in Ex 01) using next()
    # pass [Q1] & [Q2] as Q1 & Q2 arguments of the data generator.
Set batch size as 1
    # Hint: use `vocab['<PAD>']` for the `pad` argument of the data
generator
    Q1, Q2 = next(data_generator([Q1], [Q2], 1, vocab['<PAD>']))
    # Call the model
    v1, v2 = model((Q1, Q2))
    # take dot product to compute cos similarity of each pair of en
tries, v1, v2
    # don't forget to transpose the second argument
    d = np.dot(v1[0], v2[0].T)
    # is d greater than the threshold?
    res = d > threshold

    ### END CODE HERE ###

    if(verbose):
        print("Q1 = ", Q1, "\nQ2 = ", Q2)
        print("d = ", d)
        print("res = ", res)

    return res

```



```
In [44]: # Feel free to try with your own questions
question1 = "When will I see you?"
question2 = "When can I see you again?"
# 1 means it is duplicated, 0 otherwise
predict(question1 , question2, 0.7, model, vocab, verbose = True)

Q1 = [[585  76   4  46  53  21   1   1]]
Q2 = [[ 585   33   4  46   53 7280   21   1]]
d   = 0.88113236
res = True
```

Out[44]: True

Expected Output

If input is:

```
question1 = "When will I see you?"
question2 = "When can I see you again?"
```

Output is (d may vary a bit):

```
Q1 = [[585  76   4  46  53  21   1   1]]
Q2 = [[ 585   33   4  46   53 7280   21   1]]
d   = 0.88113236
res = True
True
```

```
In [45]: # Feel free to try with your own questions
question1 = "Do they enjoy eating the dessert?"
question2 = "Do they like hiking in the desert?"
# 1 means it is duplicated, 0 otherwise
predict(question1 , question2, 0.7, model, vocab, verbose=True)

Q1 = [[ 443  1145  3159  1169   78 29017   21   1]]
Q2 = [[ 443  1145   60 15302   28   78  7431   21]]
d   = 0.477536
res = False
```

Out[45]: False

Expected output

If input is:

```
question1 = "Do they enjoy eating the dessert?"
question2 = "Do they like hiking in the desert?"
```

Output (d may vary a bit):

```
Q1  = [[ 443  1145  3159  1169      78 29017      21      1]]
Q2  = [[ 443  1145      60 15302      28      78  7431      21]]
d    = 0.477536
res  = False
False
```

You can see that the Siamese network is capable of catching complicated structures. Concretely it can identify question duplicates although the questions do not have many words in common.

On Siamese networks

Siamese networks are important and useful. Many times there are several questions that are already asked in quora, or other platforms and you can use Siamese networks to avoid question duplicates.

Congratulations, you have now built a powerful system that can recognize question duplicates. In the next course we will use transformers for machine translation, summarization, question answering, and chatbots.